

# Use of Generative Adversarial Networks (GANs) and LoRA Stable Diffusion for Image Generation

**Group 3:** Matthew Jamison, James A. Nguyen,

Lawrence Mei



# What is a GAN?

GANs are models consisting of two networks:

- i. Generator
- ii. Discriminator

Role of Generator:

Fool the discriminator

Role of Discriminator:

Correctly identify fake image given one real and one generated input

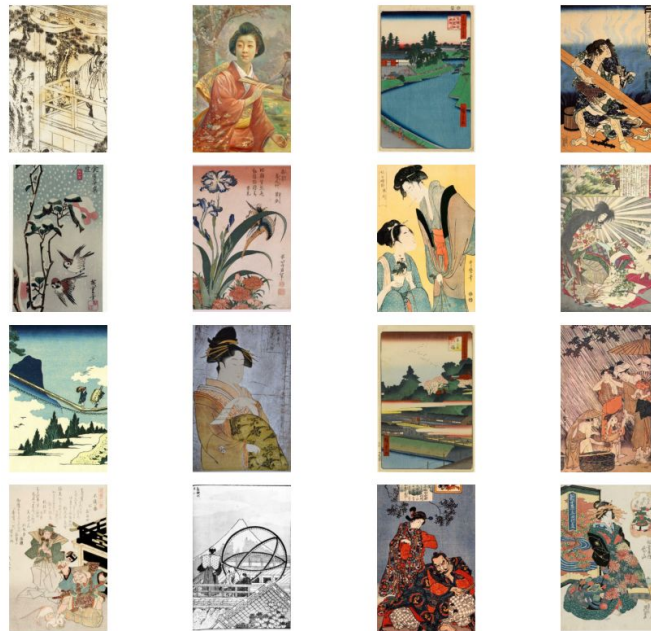
# Our Data

2236 -> 1492

Standardized resolution to 480 x 320

Manually isolated images in the style of ukiyo-e edo period style art

Idea: iterate model over random noise until resembles data

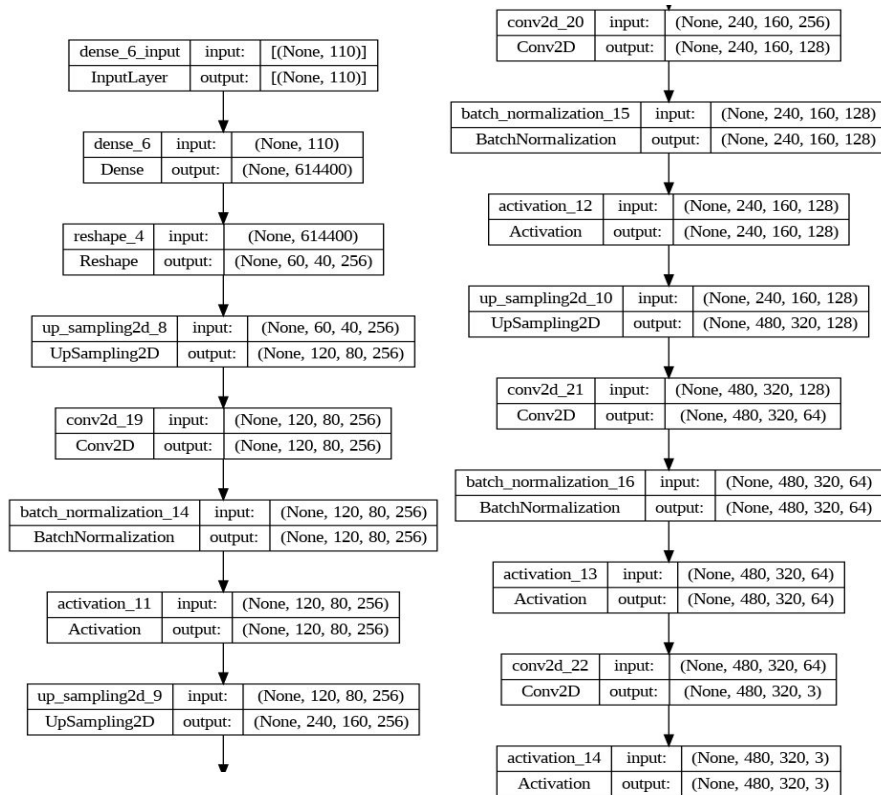


~ 1500 image subset of Edo  
Period Japanese Art

# Our Goal

- Generate images that resemble certain art styles
  - For example: ukiyo-e edo period art, cartoon, caricature, realistic, etc.
- Art restoration
- Create concept art with inspiration from artists

# Generator Network Architecture



## General Framework:

Upsampling and Batch Normalization in each layer.

ReLU activation function to avoid distortion

Decrease number of neurons with each layer

RMSprop optimizer

# Discriminator Network Architecture

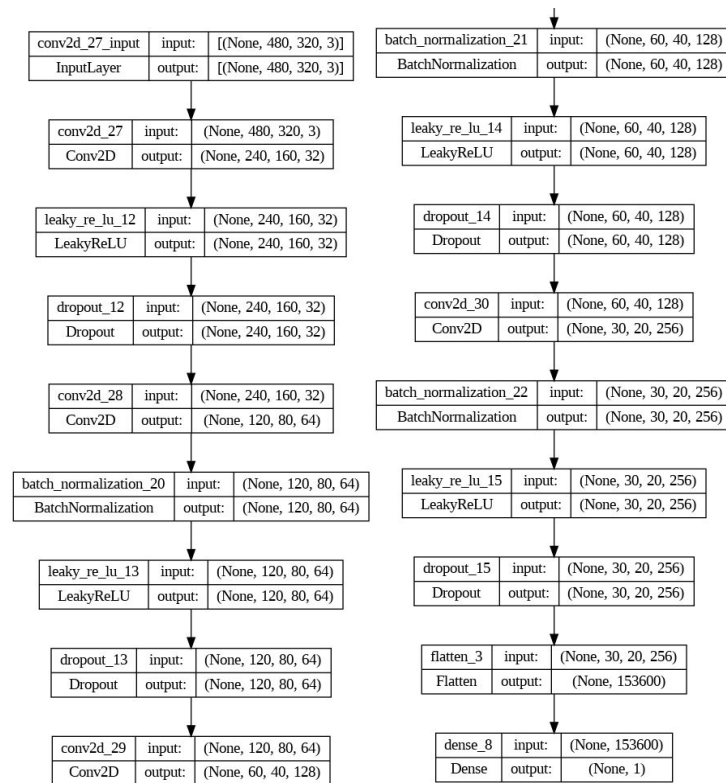
## General Framework:

Downsampling and Batch Normalization in each layer.

Leaky ReLU activation function

Increase number of neurons with each layer

RMSprop optimizer

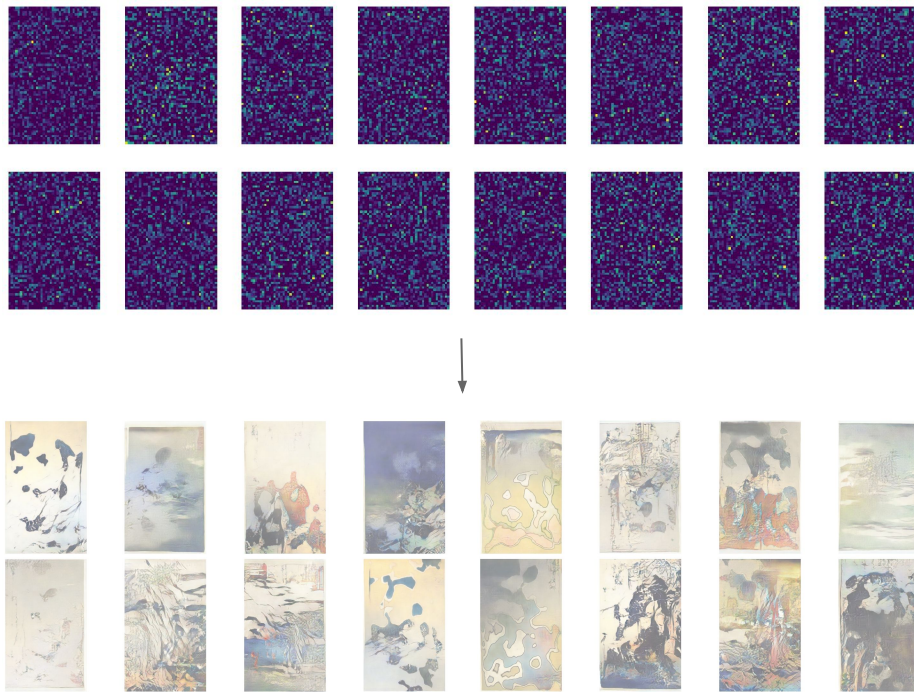


# GAN Training Loop

Fundamentally a game theory simulation

Important to balance the learning rates of the models

- i. Often involves intentionally nerfing the discriminator to even the playing field
- ii. Utilized exponential decay rate

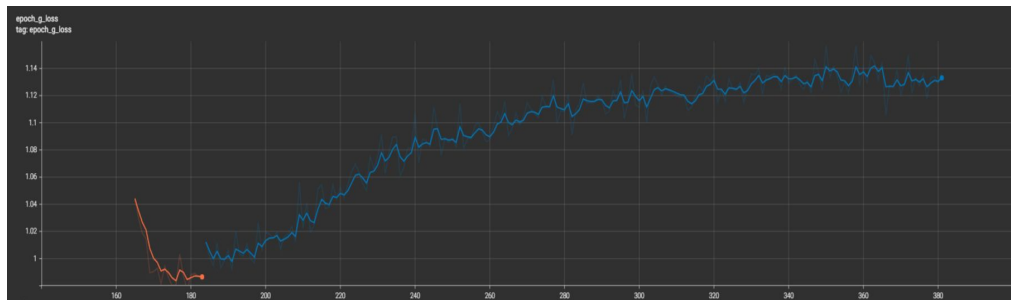
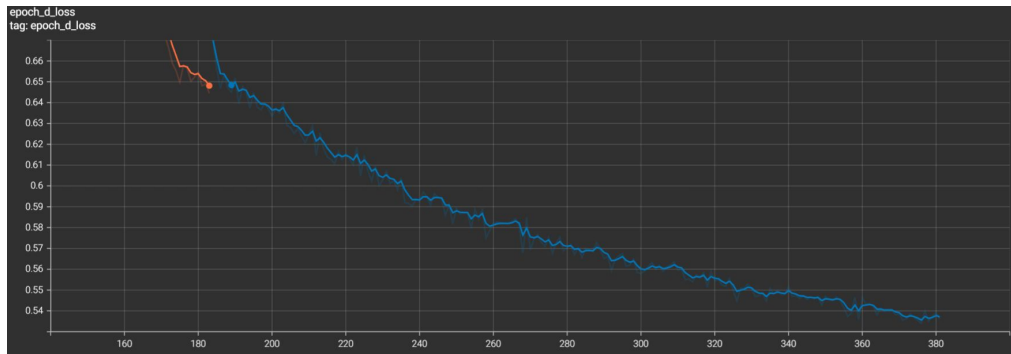


# Model Analysis

Only trained on 450 epochs before running out of compute units.

Loss chart suggests that the generator model falls behind after ~180 epochs

Could improve with larger dataset and more epochs





# Hyperparameter Tuning

Optimizer	RMSprop
Learning Rate	1e-4
Decay Steps	1000
Decay Rate	0.95
Batch Size	16
Epochs	450

RMSprop vs Adam

Equal vs Different Learning Rates

- Early Stage Mode Collapse
- Used Exponential Learning Rate Decay
  - Switch to more adaptive learning rate scheduler

Conservative Batch Size and Epochs due to resource limitations

# Model Outputs

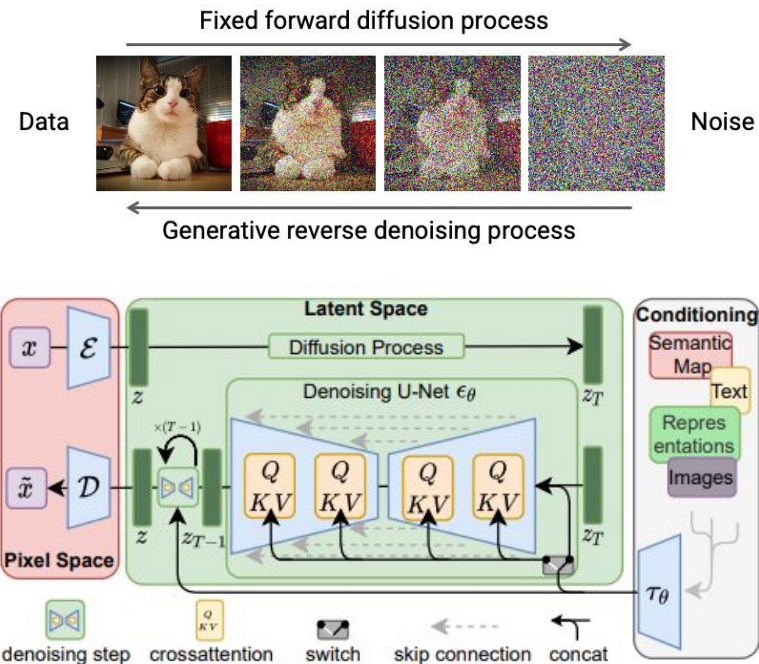
Despite training resource limitations, early training showed signs of

- i. shared qualities with the dataset e.g. color similarities
- ii. as well as simple feature extraction e.g. red tag commonly found in the top right corner



# What is Stable Diffusion?

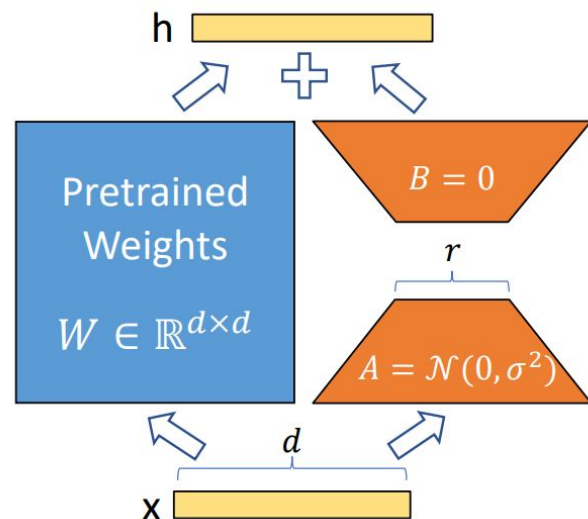
- We chose SDXL 1.0
- It is a latent diffusion model trained on 2.3 billion images
- Architecture:
  - Encoder/Decoder
  - Diffusion Process
  - Language Transformer
  - U-Net
  - Cross-Attention



# What is LoRA?

## Low-Rank Adaptation

- Pretrained model weights are frozen
- Fine-tuning weights created through low-rank decomposition to matrices A and B
- Way less training time!



# Data Preparation

- Resized images to 600x400 (GAN was 480x320)
- BLIP model to provide captions for images with some manual tweaking
- Rented GPU time on an RTX A5000 due to memory implications



a cat sitting on a window ledge looking out at a landscape

# LoRA Hyperparameter Tuning

Optimizer	Adafactor
Learning Rate	1e-3
Network Rank	128
Batch Size	20
Epochs	12

- Adafactor chosen to save on memory
- Faster learning rate due to LoRA training times
- 128 is considered “good enough” for SDXL LoRA training
- 12 epochs to compare images



# Epoch Comparison



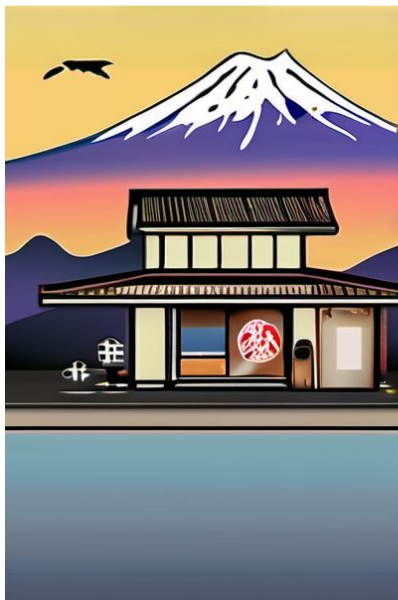
Epoch 7



Prompt: Samurai on a hill with his sword drawn, highly detailed, clouds and ocean

# Comparison to Base Model

Base Model



<lora:jpn\_art\_style-000001:1> <lora:jpn\_art\_style-000006:1> <lora:jpn\_art\_style-000007:1>



Prompt: Lawson's store with Mt. Fuji in the background



# More Results



# Future Works and Improvements

- Wasserstein loss
  - Relationship between  $y_{\text{true}}$  and  $y_{\text{false}}$
- Indicators such as frechet inception distance (FID) and inceptions score to evaluate generator quality
- More training data, more epochs
- Higher network rank for LoRA training

Thank you.

